

PAR Laboratory Assignment
Lab 3: Iterative task decomposition with OpenMP:
the computation of the Mandelbrot set

E. Ayguadé, J. R. Herrero, P. Martínez-Ferrer,
J. Morillo, J. Tubella and G. Utrera

Spring 2021-22

Index

Index	1
1 Before starting this laboratory assignment ...	2
1.1 Iterative task decompositions	2
1.2 Should I remember something from previous laboratory assignments?	3
1.3 So, what should I do next?	3
2 Task decomposition analysis for the Mandelbrot set computation	4
2.1 The Mandelbrot set	4
2.2 Task decomposition analysis with <i>Tareador</i>	5
3 Implementing task decompositions in <i>OpenMP</i>	7
3.1 <i>Point</i> strategy implementation using task	7
3.2 <i>Point</i> strategy with granularity control using taskloop	9
3.3 <i>Row</i> strategy implementation	10
3.4 Optional	10
4 Deliverable	11
4.1 Task decomposition and granularity analysis	11
4.2 <i>Point</i> decomposition in <i>OpenMP</i>	11
4.3 <i>Row</i> decomposition in <i>OpenMP</i>	12
4.4 Optional	12

Note:

- This laboratory assignment will be done in two sessions (2 hours each). For the first session we suggest you do chapter 2 and start with chapter 3.
- All files necessary to do this laboratory assignment are available in a compressed tar file available from the following location: `/scratch/nas/1/par0/sessions/lab3.tar.gz`. Copy it to your home directory in `boada.ac.upc.edu` and uncompress it with this command line: `"tar -zxvf lab3.tar.gz"`.

1

Before starting this laboratory assignment ...

Before going to the classroom to start this laboratory assignment, we strongly recommend that you take a look at this section and try to solve the simple questions we propose to you. This will help to better face your first programming assignment in OpenMP: the Mandelbrot set computation.

1.1 Iterative task decompositions

As you have already realised, loops and other iterative constructs are one of the main sources of parallelism in programs. **Iterative task decompositions** are parallelisation strategies that try to exploit this parallelism. For example, consider the simple loops in Figure 1.1 computing a) the elements of vector **C** as the sum of the elements of vectors **A** and **B** and b) the dot product of vectors **A** and **B**.

```
void vectoradd(int *A, int *B, int *C, int n) {  
    for (int i=0; i< n; i++)  
        C[i] = A[i] + B[i];  
}
```

a)

```
int dotproduct(int *A, int *B, int n) {  
    int sum=0;  
    for (int i=0; i< n; i++)  
        sum += A[i] * B[i];  
    return(sum);  
}
```

b)

Figure 1.1: Simple examples performing a) vector sum and b) dot product.

In an iterative task decomposition tasks correspond with the execution of one or more loop iterations. The granularity of the task decomposition would be determined by the number of iterations executed per task; iterations do not need to be consecutive, but sometimes this may help to better exploit cache locality and avoid other artefacts you will be learning during the course.

An iterative task decomposition is named "embarrassingly parallel" if the execution of all tasks can be performed totally in parallel without the need of satisfying data sharing and/or task ordering constraints. This is the case in the example in the upper part of Figure 1.1; there are no data races preventing the execution of all iterations in parallel. However the computation in the example in the lower part of the same figure will require some sort of synchronisation in order to avoid the possible data races caused by the access to the variable **sum**.

1.2 Should I remember something from previous laboratory assignments?

Would you be able to write a parallel version in *OpenMP* for these two simple code fragments? We are quite sure that your answer is "Yes!, of course". In the previous laboratory assignment you should have learned about the different options in *OpenMP*, and when to use them, to express tasks out of loops (either with implicit tasks or explicit tasks with `task` and `taskloop`), with the appropriate thread creation (`parallel` and `single`) and how to enforce task order with task barriers (`taskwait` and `taskgroup`), and data sharing constraints (`critical`, `atomic` and `reduction` operations).

In addition you should remember from the first assignment how to use the *Tareador* API and GUI to understand the potential parallelism available in a sequential code, as well as the causes that limit this parallelism. And also the use of *Extrae* and *Paraver* to visualise the execution of your parallel *OpenMP* program and understand its performance.

1.3 So, what should I do next?

Simply we ask you to think about, better to write in paper, the multiple alternatives to code the parallel versions in *OpenMP* for the two simple codes shown in Figure 1.1. You don't need to deliver them, but we will comment your different solutions in the lab session, if necessary.

2

Task decomposition analysis for the Mandelbrot set computation

2.1 The Mandelbrot set

In this laboratory assignment you are going to explore the tasking model in *OpenMP* to express **iterative task decompositions**. But before that you will start by exploring the most appropriate ones by using *Tareador*. The program that will be used is the computation of the *Mandelbrot set*, a particular set of points, in the complex domain, whose boundary generates a distinctive and easily recognisable two-dimensional fractal shape (Figure 2.1).




Figure 2.1: Fractal shape

For each point c in a delimited two-dimensional space, the complex quadratic polynomial recurrence $z_{n+1} = z_n^2 + c$ is iteratively applied n to determine if it belongs or not to the Mandelbrot set. The point is part of the Mandelbrot set if, when starting with $z_0 = 0$ and applying the iteration repeatedly, the absolute value of z_n never exceeds a certain number however large n gets.

A plot of the Mandelbrot set is created by colouring each point c in the complex plane with the number of steps max for which $|z_{max}| \geq 2$ (or simply $|z_{max}|^2 \geq 2 * 2$ to avoid the computation of the square root in the modulus of a complex number). In order to make the problem doable, the maximum number of steps is also limited: if that number of steps is reached, then the point c is said to belong to the Mandelbrot set.

If you want to know more about the Mandelbrot set, we recommend that you take a look at the following Wikipedia page:

http://en.wikipedia.org/wiki/Mandelbrot_set

In the *Computer drawings* section of that page you will find different ways to render the Mandelbrot set.

1. Open the `mandel-seq.c` sequential code to identify the loops that traverse the two dimensional space (loops with induction variables `row` and `column` and the loop that is used to determine if a point belongs to the Mandelbrot set (`do--while` loop).
2. Compile the sequential version of the program by using `"make mandel-seq"`. You will get a binary that can be executed with different output options: 1) the Mandelbrot set is computed but no output is displayed or written to disk (just for timing purposes of the computational part); 2) `-h` option: the histogram for the values in the Mandelbrot set is also computed; 3) `-d` option: the Mandelbrot set is displayed, for visual inspection; and 4) `-o` option: the values for Mandelbrot set and/or histogram are written to disk, to compare the numerical values obtained with the reference output. When executed with the `-d` option the program will open an X window to draw the set; once the program draws the set, it will wait for a key to be pressed; in the meanwhile, you can find new coordinates in the complex space by just clicking with the mouse (these coordinates could be used to zoom the exploration of the Mandelbrot set).
3. The program includes additional parameters to specify the domain in the complex space where the computation of the Mandelbrot set has to be performed. Execute `"./mandel-seq -help"` to see the options that are available to run the program. For example the `"-i"` specifies the maximum number of iterations at each point (default 1000) and `-c` and `-s` specify the center $x_0 + iy_0$ of the square to compute (default origin) and the size of the square to compute (default 2, i.e. size 4 by 4). For example you can try `"./mandel-seq -d -c -0.737 0.207 -s 0.01 -i 100000"` to obtain a different view of the set.
4. In order to have a reference of the sequential execution, execute the sequential version with `"./mandel-seq -h -i 10000 -o"` to get its execution time and the output file, to be used later to check the correctness of the different parallel versions you will write.

As you can guess at this point, the computation of the Mandelbrot set problem is totally amenable for applying the iterative task decomposition parallelisation strategy. Figure 2.2 shows the double-nested loop that you have found in the source code for the sequential version in `mandel-seq.c`.

```
for (int row = 0; row < height; ++row) {
    for (int col = 0; col < width; ++col) {
        ... // computation of a single point in the Mandelbrot set
    }
}
```

Figure 2.2: Two loops traversing the iteration space for the Mandelbrot set.

Two different alternatives to generate tasks will be explored in this laboratory assignment, with different granularities: a) *Row* decomposition in which a task will correspond with the computation of one or more (consecutive) rows of the Mandelbrot set; and b) *Point* decomposition in which a task will correspond with the computation of one or more (consecutive) points in the same row of the Mandelbrot set. In other words, *Row* will correspond with the distribution of iterations of the loop indexed by the `row` induction variable while *Point* will correspond with the distribution of iterations of the loop indexed by the `col` induction variable, for a fixed value of `row`. The granularity of both *Row* and *Point* decompositions will be determined by the number of iterations per task in the respective loops.

2.2 Task decomposition analysis with *Tareador*

In this section you will study the main characteristics of each proposed task decomposition strategy (*Row* and *Point*) using *Tareador*, in both cases with a granularity of one iteration per task. Then, in the next session you will implement them using the *OpenMP* programming model and explore different task granularities. We recommend that you follow the guidelines in the following bullets and reach the appropriate conclusions.

1. Complete the sequential `mandel-tar.c` code partially instrumented in order to analyse the potential parallelism for the *Row* strategy, with granularity of one iteration of the `row` loop per task. Then

compile the instrumented code using the appropriate target in the `Makefile` that generate the binary for analysis with *Tareador*.

2. Interactively execute `mandel-tar` binary using the `./run-tareador.sh` script, only indicating the name of the instrumented binary (with no additional options). The size of the image to compute is defined inside the script (we are using `-w 8` as the size for the Mandelbrot image in order to generate a reasonable task graph in a reasonable execution time). Which are the two most important characteristics for the task graph that is generated? Save the TDG generated for later inclusion in the deliverable.
3. Next interactively execute `mandel-tar` binary using the `./run-tareador.sh`, but now indicating the name of the instrumented binary and the `-d` option. **Important:** look for a small window that will be opened to display the tiny Mandelbrot set that is computed and click inside in order to finish its execution. Again, which are the two most important characteristics for the task graph that is generated? Which part of the code is making the big difference with the previous case? How will you protect this section of code in the parallel *OpenMP* code that you will program in the next sessions? Save the new TDG generated for later inclusion in the deliverable.
4. Finally, interactively execute `mandel-tar` binary using the `./run-tareador.sh`, but now indicating the name of the instrumented binary and only the `-h` option. What does each chain of tasks in the task graph represents? Which part of the code is making the big difference with the two previous cases? How will you protect this section of code in the parallel *OpenMP* code that you will program in the next sessions? Don't forget to save the new TDG generated for later inclusion in the deliverable.

Once the study for the *Row* strategy is completed, modify the instrumented `mandel-tar.c` code to analyse the potential parallelism for the *Point* strategy. Repeat the analysis for the three previous executions. Are conclusions the same? Which is the main change that you observe with respect to the *Row* strategy?

Important: Please, refer to section 4.1 to make sure you have everything you need to include in the deliverable for this laboratory assignment.

3

Implementing task decompositions in *OpenMP*

In this session you will explore different options in the *OpenMP* tasking model to express the iterative task decomposition strategies for the Mandelbrot computation program. You will analyse the scalability and behaviour of your implementation using the instrumentation and analysis tools you should start to be familiar with.

3.1 *Point* strategy implementation using task

The simplest way to create a task in *OpenMP* for the computation of each point in the Mandelbrot set (*Point* task decomposition) is shown in Figure 3.1: each iteration of the `col` loop will be executed as an independent task.

```
#pragma omp parallel
#pragma omp single
for (int row = 0; row < height; ++row) {
    for (int col = 0; col < width; ++col) {
        #pragma omp task firstprivate(row, col)
        {
            ...
        }
    }
}
```

Figure 3.1: Simplest tasking code for *Point* granularity

But before that, the program needs to create a team of threads that will execute the tasks generated by the previous pragma. The team of threads is created once by using the `parallel` construct; immediately after that, a single task generator is specified by using the `#pragma omp single` work-distributor in *OpenMP*: the thread that gets access to the `single` region, traverses all iterations of the `row` and `col` loops, generating a task for each iteration of the innermost loop; the rest of threads (i.e. those that do not gain access to the `single`) simply wait at the implicit barrier at its end and, in the meanwhile, execute any of the tasks that are generated. As indicated in the `task` pragma, each task gets a private copy of variables `row` and `col` initialised with the value they have at task creation time (`firstprivate` clause). With this, when one of the tasks is extracted from the pool by any of the threads in the team, it has the values for these two variables necessary for the computation of a specific point in the Mandelbrot set. Once all tasks are finished, threads will leave the implicit barriers at the end of the `single` and `parallel` and one of them continue with the execution of the sequential part.

1. Edit the initial task version in `mandel-omp.c`. Check that it corresponds with the *Point* parallelization shown in Figure 3.1 and **insert the missing *OpenMP* directives to make sure that all the dependences you detected with *Tareador* in the previous section are honoured**. Make sure you use **the most efficient mechanism** to protect each dependence. You should not worry about the two "fake" parallel regions that have been added to delimit its start and end of the `main` program; this is used by the scripts that will trace the execution with *Extræ*, cutting the trace appropriately for your proper analysis.
2. Use the `mandel-omp` target in the `Makefile` to compile and generate the parallel binary for this first version of the parallel code. In order to visually check the correctness of your parallelization, interactively execute it with the `-d` option to see what happens when running it with only 1 thread and a maximum of 10000 iterations per point (i.e. `OMP_NUM_THREADS=1 ./mandel-omp -d -h -i 10000`). Is the image generated correct? Now interactively execute the binary with 2 threads and the same number of iterations per point (e.g. `OMP_NUM_THREADS=2 ./mandel-omp -d -h -i 10000`). Is the image generated still correct?
3. In order to measure the reduction in time due to the parallel execution, submit the execution of the `submit-omp.sh` script indicating the `mandel-omp` binary with 1 and 8 threads. This script does not include the `-d` option but generates an output file using the `-o` option. At this point it is important that you verify that the output files that are generated are correct, by comparing them with the output file obtained from the original sequential version (if necessary compile the binary and execute it). Take a look at the execution times that are reported after the two executions. Is the speed-up appropriate? The better answer this question and have a broader picture, submit the `submit-strong-omp.sh` script with `sbatch` to execute the binary generated and obtain the execution time and speed-up plots for the 1–12 processor range. Visualise the *PostScript* file generated. Again, is the scalability appropriate?
4. In order to better understand how the execution goes, submit an *Extræ* instrumented execution with 8 threads using the `submit-extræ.sh` script, opening the trace generated with *Paraver*. The instrumented execution uses a smaller Mandelbrot set (320x320, set with the `-w 320`, **without changing the maximum number of iterations `-i 10000` option inside the `submit-extræ.sh` script**) in order to reduce the size of the trace and the time it takes to generate and process it.
 - (a) Spend a couple of minutes to remember how the parallel execution works. Unfold the *Hints* menu in the main *Paraver* window and visualise the *implicit tasks* originated in the `parallel` construct and how they enter into the `single` work-sharing construct. Only one implicit task gains the access to it and stays there until the end of the parallel region. You can also visualise when the implicit tasks fall into the implicit barrier synchronisation at the end of `single`, sitting there waiting for explicit tasks to be executed.
 - (b) Next look for the appropriate configuration file to visualise when *explicit tasks* are created and executed. Which threads are creating and executing them? How many tasks are created/executed? For this, look for the appropriate configuration file in the *Hints* menu to obtain a profile of the tasks generated and executed. Two profiles are displayed, one for the executed tasks and another for the created tasks; in the *Statistic* option in the control window you can select the metric you want to visualize: the default metric is the number of tasks but you can change to see, for example, the average time to execute a task or to create a task, the total time spent in executing/creating tasks or the standard deviation. Is the load well balanced (both in terms of number of tasks and in terms of total time executing tasks)? Which of the two is actually relevant?
 - (c) Find the appropriate configuration file in the *Hints* menu that opens the histogram showing the duration of explicit tasks. Each bin of the histogram shows the number of tasks executed by the thread with a specific task duration. Observe the high dispersion in task durations (look for the duration of the shortest and longest tasks). You can change the *Statistic* to see for example % of time executing tasks with a certain task duration. For this histogram the *3D.plane* selector shows the only task type that is generated and executed in our program.

- (d) Finally look at the overheads related with synchronization and task scheduling (which hint in the list can provide you with this information?). In summary, is granularity of the tasks appropriate? Would it be easy to change?

Important: Take note of all the relevant information in order to draw some conclusions about the behaviour of this code version. You will have to include them in the deliverable for this laboratory assignment and support with the appropriate *Paraver* window captures and performance plots.

3.2 *Point* strategy with granularity control using taskloop

Next you will make use of the `taskloop` construct, which generates tasks out of the iterations of a `for` loop, allowing to better control the number of tasks generated or their granularity (i.e. the number of iterations per task). As you know, you can make use of the `num_tasks` clause to control the number of tasks generated out of the loop, with the appropriate number to specify the number of tasks; alternatively, you can use the `grainsize` clause to control the granularity of tasks, with the appropriate number to specify the number of iterations per task. If none of these clauses is specified, the *OpenMP* implementation decides a value for them, depending on the number of threads in the parallel region.

1. Edit the last version in `mandel-omp.c` in order to make use of `taskloop` to implement the *Point* strategy. At this point you can leave unspecified the granularity for the tasks (i.e. do not use neither `num_tasks` nor `grainsize`), leaving to the *OpenMP* implementation to take the decision about the granularity to apply.
2. Compile this new version and interactively execute it with 1 and 2 processors with the `-d` option to verify that the code still visualizes the correct Mandelbrot set. Also submit the execution of the `submit-omp.sh` script for 1 and 8 processors to validate the output files generated with respect to the output of the original sequential program and to observe the new execution times. Is the speed-up appropriate? In order to better answer to this last question, submit the `submit-strong-omp.sh` script and visualise the new scalability plot that is obtained. Is the version with `taskloop` performing better than the last version based on the use of `task`?
3. Trace the execution of the new binary with 8 processors in order to better understand the scalability obtained. Opening the trace generated with *Paraver* and using the same configuration files as before, answering again the same questions that we posed (point 4 in the previous section 3.2). Observe the granularity for the tasks that are executed and the number of tasks that are instantiated (one per `taskloop`) vs. the number of tasks that are executed. How many tasks are executed per `taskloop`? Does this number depends on the number of threads that are used to execute?
4. Look for the appropriate hint to visualise which kind of task synchronisation (`taskwait` or `taskgroup`) is being used in this version of the program. Where are these task synchronisations happening? Observe that the execution does not continue until all tasks generated up to each one of these points are finished, introducing a certain load unbalance. Do you think these task barriers are necessary? Justify your answer. If your answer was negative, you can add the `nogroup` clause to eliminate the implicit `taskgroup`. Edit the code and submit again the `submit-strong-omp.sh` script. Has the scalability improved?. Trace again the execution with 8 threads and observe how tasks are now executed.

Important: Please, refer to section 4.2 to make sure you have everything you need (relevant information, appropriate *Paraver* captures and performance plots) to draw conclusions out of the analysis performed for the *Point* strategy in order to include them in the deliverable for this laboratory assignment.

3.3 *Row* strategy implementation

Based on the experience and conclusions you obtained from the previous sections, finally we ask you to write a new parallel version for `mandel-omp.c` that obeys to the *Row* strategy.

1. Check that the result is correct and do the analysis of its strong scalability.
2. Trace the execution for 8 processors to conclude about the performance that is obtained with this coarser-grain decomposition. Use again all previous configuration files to support your reasoning and draw the attention to the main differences that you observe when comparing the *Row* and *Point* granularities. Analyse the number of tasks created vs. the number of tasks executed and their granularities. Is there any load unbalance? Would the use of `grainsize` reduce it? Any task synchronisation to eliminate?

Important: Please, refer to section 4.3 to make sure you have everything you need to draw conclusions out of the analysis performed for the *Row* strategy in order to include them in the deliverable for this laboratory assignment.

3.4 Optional

Explore how the best versions for *Point* and *Row* behave in terms of performance for different task granularities, i.e. when setting the number of tasks or the number of iterations per task to a different value than the one chosen by *OpenMP*. To do that we provide the `submit-numtasks-omp.sh` script, which explores different granularities by changing the number of tasks that are executed in each `taskloop` (800, 512, 256, 128, 64, 32, 16, 8, 4, 2 and 1 tasks) for a certain number of threads (which is the only argument of the script). But you should first investigate how the script is passing this value to the program, and then modify the `mandel-omp.c` to receive it and use in the `num_tasks` clause of the `taskloop` construct. The script plots the results in a graphical way; reason about the results that are obtained.

Important: Please, refer to section 4.4 to make sure you have everything you need to include your conclusions in the deliverable for this laboratory assignment. j

4

Deliverable

You have to write a short document that presents the main results and conclusions that you have obtained when doing this assignment. Only PDF format for this document will be accepted.

- The document should have an appropriate structure, including, at least, the following sections: Introduction, Task decomposition strategies, Implementation in OpenMP and performance analysis, and Conclusions. The document should also include a front cover (assignment title, course, semester, students names, the identifier of the group, date, ... and a table of contents).
- **Important:** You DON'T have to include in the document all the steps you have followed during the laboratory sessions, ONLY the main results and conclusions derived from them. In the following subsections we highlight the aspects that you should included in the document. Your explanations should be based on the relevant execution timelines and other *Paraver* windows as well as scalability plots generated by the execution scripts.
- You also have to deliver the complete C source codes for Tareador instrumentation and all the OpenMP parallelisation strategies that you have done. Your professor should be able to re-execute the parallel codes based on the files you deliver. DON'T include code in the document unless strictly necessary, and in that case, only the fragment modified with respect to the original one.

Your professor will open the assignment in *Atenea* and set the appropriate dates for the delivery. You will have to deliver TWO files, one with the document in PDF format and one compressed file (*tgz*, *.gz* or *.zip*) with the requested source codes.

As you know, this course contributes to the **transversal competence "Tercera llengua"**. Deliver your material in English if you want this competence to be evaluated. Please refer to the "Rubrics for the third language competence evaluation" document to know the *Rubric* that will be used.

4.1 Task decomposition and granularity analysis

Explain the different task decomposition strategies and granularities explored with *Tareador* for the *Mandelbrot* code, including in your deliverable the task dependence graphs obtained when using the different execution options: *-d* option, *-h* option or none of both. Clearly indicate the most important characteristics for the graphs that you obtained. Explain which section of the code is causing the serialisation of tasks when using the *-d* and the *-h* options and how this section of code should be protected when parallelising with *OpenMP*. Reason when each strategy/granularity should be used.

4.2 *Point* decomposition in *OpenMP*

For the *Point* strategy implemented in *OpenMP*, describe and reason about how the performance has evolved for the three versions of the code (one with *task* and 2 with *taskloop*) that you have evaluated, using the speed-up plots obtained in the strong scalability analysis and the relevant set of *Paraver* captures to justify the performance differences you have observed.

4.3 *Row* decomposition in *OpenMP*

For the *Row* strategy implemented in *OpenMP*, describe the parallelization strategy that you have decided to implement. Reason about the performance that is obtained comparing with the results obtained for the best *Point* implementation, including the execution time and speed-up plots obtained in the strong scalability analysis and the relevant *Paraver* captures that help you to better explain the results that have been obtained.

4.4 Optional

If you decided to analyse the influence of task granularity in both *Point* and *Row* strategies, report the conclusions that you reached, basing them on the performance plots obtained and *Paraver* windows that may support your explanations. Briefly explain how did you changed the code in `mandel-omp.c` to interact with the exploration script `submit-numtasks-omp.sh` that we provided.